

Mixture distributions in collaborative probabilistic forecasting of disease outbreaks

by

Spencer Gordon Wadsworth

A Creative Component submitted to the graduate faculty
in partial fulfillment of the requirements for the degree of
MASTER OF SCIENCE

Major: Statistics

Program of Study Committee:
Jarad Niemi, Major Professor
Karin Dorman
Kori Khan

Iowa State University

Ames, Iowa

2022

Copyright © Spencer Gordon Wadsworth, 2022. All rights reserved.

DEDICATION

I would like to dedicate this manuscript to my uncle Bryce. At his recommendation I considered writing this about him, but then I decided I wanted to graduate.

TABLE OF CONTENTS

	Page
LIST OF TABLES	iv
LIST OF FIGURES	v
ACKNOWLEDGMENTS	vi
ABSTRACT	vii
1 Introduction	1
1.1 CDC Influenza Like Illness	1
1.2 COVID-19 Forecast Hub	2
1.3	3
2 Probabilistic Forecast Representations	5
2.1 Representation properties to consider in collaborative projects	5
2.2 Probabilistic Forecast Representation Types	6
2.3 Discrete mixture of parametric distributions	17
3 Mixture distributions in a collaborative forecast project	21
3.1 Submission format	21
3.2 Mixture construction and scoring tools	22
3.3 Ensemble construction	25
4 Retrospective Analysis	27
4.1 CDC Influenza like Illness	27
4.2 COVID-19 Forecast Hub	31
5 Discussion	36
6 References	36
7 Appendix	39

LIST OF TABLES

		Page
1	Parametric distribution storage	9
2	Discretized bin distribution storage	14
3	Quantile/interval and values storage	17
4	Forecast representation comparison	20
5	Influenza competition submission example	21
6	Mixture distribution forecast submission example	22
7	Toy prediction 1	23
8	Toy prediction 2	23
9	Discretized bin distribution cutoff values	29
10	CDC influenza retro analysis results	31
11	Quantile cutoff values	32
12	COVID-19 Forecast Hub results	33

LIST OF FIGURES

		Page
1	Official CDC COVID-19 deaths report August 2021	4
2	Density/CDF comparison between parametric distribution, sample distribution, discretized bin distribution and quantiles	18
3	Toy mixture distributions	24
4	Toy ensemble forecast	27
5	Parametric distribution fits to binned distribution	30
6	QQ plot for quantile fit	33
7	CDF plot for quantile fit	34

ACKNOWLEDGMENTS

ABSTRACT

Collaboration among multiple teams has played a major role in probabilistic forecasting of influenza outbreaks, the COVID-19 pandemic, other disease outbreaks and problems in many other fields. When collecting forecasts from individual teams, ensuring that each team's model represents forecast uncertainty according to the same format allows for direct comparison of forecasts as well as methods of constructing multi-model ensemble forecasts. This paper outlines several common probabilistic forecast formats including probability densities, samples, binned distributions and quantile forecasts and compares their usage within the context of collaborative projects. We propose the use of a discrete mixture distribution format in collaborative forecasting in place of other formats. The flexibility in distribution shape, the ease for scoring and building ensemble models and the reasonably low level of computer storage required to save such a format make the discrete mixture distribution an attractive alternative to the other mentioned formats.

1 Introduction

Predicting the outcomes of prospective events is the object of much scientific inquiry and the basis for many decisions both public and private. Because predictions of the future can almost never be precise, it is usually desirable that a level of uncertainty be attached to any prediction. In recent years, it has become increasingly desirable that forecasts be probabilistic in order to account for uncertainty in predicted quantities or events [Gneiting and Katzfuss \(2014\)](#). Specific problems for which probabilistic forecasting is used include weather forecasting [Baran and Lerch \(2018\)](#), economics [Groen et al. \(2013\)](#) and disease outbreaks [Yamana et al. \(2016\)](#).

A probabilistic forecast is a forecast to which probabilities are assigned to various possible outcomes. There are a number of ways whereby probabilities or uncertainty may be represented. A common representation is either a continuous or discrete parametric distribution, given as a probability density/mass function (pdf/pmf). Much of the literature on calibration, sharpness and scoring of a forecast pertains to parametric distribution forecasts [Gneiting et al. \(2007\)](#); [Gneiting and Ranjan \(2013\)](#); [Baran and Lerch \(2018\)](#). Other common representations include samples [Krueger et al. \(2016\)](#), discretized bins [McGowan et al. \(2019\)](#) and quantile or interval type forecasts [Taylor \(2021\)](#) [Bracher et al. \(2021\)](#). Each representation may be more or less appropriate than the others for a given problem but knowing how to interpret, score and construct ensemble forecasts for a selected representation is essential when multiple teams collaborate in the same forecasting project.

Two collaborative projects on forecasting disease outbreaks for which many separate forecasts are used include the United States Centers for Disease Control (CDC) hosted annual competition for forecasting influenza [CDC \(2021\)](#) and the COVID-19 Forecast Hub which has continuously operated since the start of the COVID-19 pandemic in the US in early 2020 [Cramer et al. \(2021a\)](#).

1.1 CDC Influenza Like Illness

Since the 2013-14 flu season, the CDC has hosted an annual competition for predicting the timing, peak and intensity of the year's flu season. Forecasts for the different targets also include

predictions for one, two, three and four weeks ahead of the prediction time. National flu data is provided weekly to outside academic teams who use that data to construct forecasts using whatever methods they choose. Historically, the forecast predictions must be submitted in a discretized bin format. During previous flu seasons the binning scheme was on a bounded numeric scale and the prediction of a specific target was a set of probabilities assigned to each bin [McGowan et al. \(2019\)](#). These forecasts are then evaluated against actual flu activity, and at the end of the season a winning team is declared [CDC \(2021\)](#).

This competition has provided the CDC, competing teams and other interested parties a chance to collaborate and improve forecasting each season. One proposed way to enhance prediction has been to aggregate the various teams' forecasts into a multi-model ensemble forecast [McGowan et al. \(2019\)](#); [McAndrew and Reich \(2019\)](#); [Reich et al. \(2019b\)](#).

A multi-model ensemble forecast is a combination of several component forecast models into one model which often yields better predicting power than the individual models [Cramer et al. \(2021b\)](#). Such an ensemble made from multiple influenza competition forecasts did in fact outperform the individual component models [Reich et al. \(2019b\)](#). In the COVID-19 Forecast Hub, construction of ensemble forecasts is a main priority.

1.2 COVID-19 Forecast Hub

In March 2020, at the onset of the COVID-19 pandemic, the COVID-19 Forecast Hub was founded. Borrowing on work and ideas from the CDC influenza competition, the Forecast Hub is a central site in which dozens of academic teams collaborate to forecast the ongoing COVID-19 pandemic. Every week relevant pandemic data is provided to these teams who build probabilistic models for their use in forecasting cases, hospitalizations and deaths due to COVID-19. Forecasts are made on the US county, state and national level with predictions for days, weeks and months ahead. These forecasts are combined into a single ensemble forecast. The model data, forecasts and the ensemble are passed along to the CDC for its use in official communication [Cramer et al. \(2021a\)](#). See Figure 1.

Though similar to the forecasting in the influenza competition, the format of the COVID-19 Forecast Hub has some key distinctions. For one this project has been operating continuously since it first began, so forecasts have been made for over 100 straight weeks. As well due to the initial lack of understanding of the COVID-19 virus and the time pressure of creating forecasts, rather than requesting forecasts as binned probabilities the forecasts are requested as the predictive median and predictive intervals for various nominal levels depending on the target to be predicted [Bracher et al. \(2021\)](#). Collecting forecasts in this quantile or interval type format brings with it differences in how to score, construct ensemble models and store the forecasts among other differences.

1.3

In the context of collaborative forecasting like that of the CDC flu competition or the COVID-19 Forecast Hub, discretized bins and quantile forecasts have proven useful and effective. Yet both types come with their drawbacks. Data storage for instance might be a concern if many bins are used for forecasting. And scoring methods are limited if forecasts are made of prediction intervals.

In this paper, we propose the use of discrete mixture distributions as a means of forecasting for projects similar to the CDC flu competition or the COVID-19 Forecast Hub. In section 2, popular probabilistic forecast representation types are defined and reviewed. Methods of scoring, storing, building ensembles and other aspects will be reviewed and compared for each representation type. In section 3, details for using discrete mixture distributions in a collaborative forecast project are presented as well as tools which may be used for scoring and ensemble building. Section 4 is a retrospective study of the CDC flu competition and COVID-19 forecast predictions and an attempt to assess whether any predictions approximately come from well known continuous distributions. This paper is concluded with a discussion.

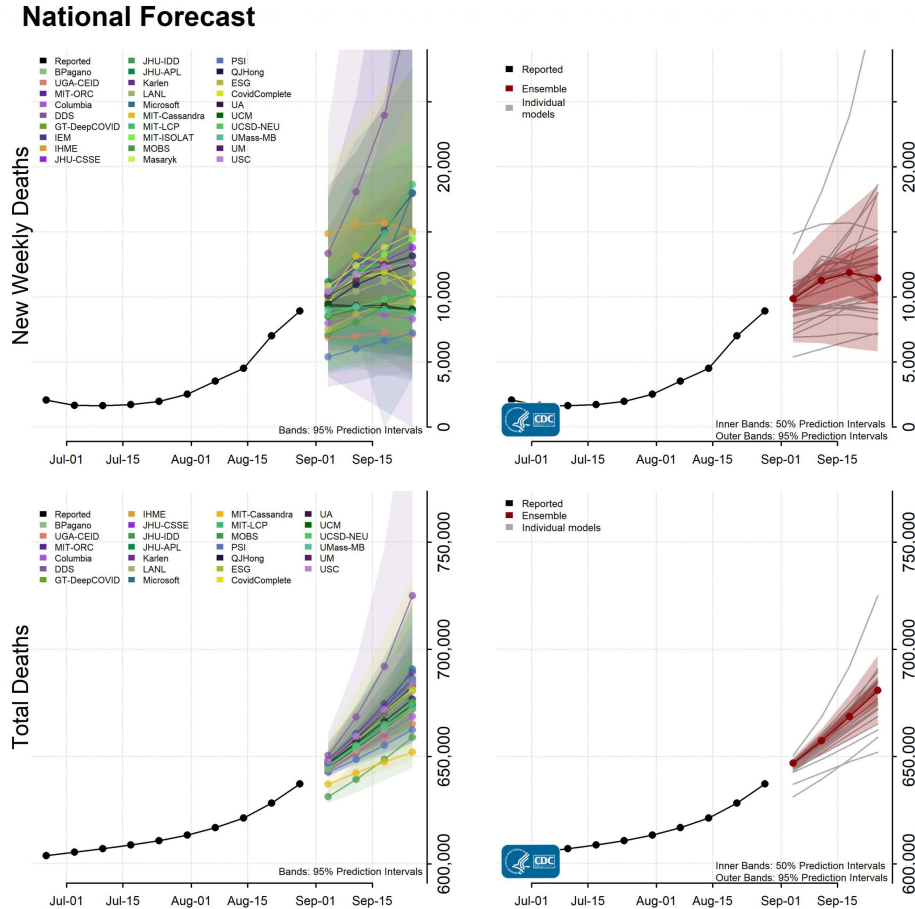


Figure 1: This image, published on [cdc.gov](https://www.cdc.gov) as official public communication by the CDC, shows forecast models for national new weekly deaths due to COVID-19 in the top row and cumulative deaths on the bottom row. The plots in the left column show prediction intervals from multiple teams whereas the plots on the right show the intervals for an ensemble forecast model.

2 Probabilistic Forecast Representations

Before introducing discrete mixture distributions as a representation type for a probabilistic forecast, we review representations already commonplace in forecasting. In a collaborative setting, certain aspects of each representation should be considered. With each representation presented, common applications of scoring, data storage and ensemble construction will be considered for as well as other notable properties.

2.1 Representation properties to consider in collaborative projects

2.1.1 Scoring

Scoring rules are used to numerically evaluate or score a probability forecast. The score is a measure of the accuracy of the forecast and where multiple forecasts exist the score for each may be used to compare forecasts. If a scoring rule is proper, then the best possible score is obtained by reporting the true distribution. The rule is strictly proper if that value is unique. Under proper scoring rules, a forecaster has no incentive to be dishonest in their submission [Gneiting and Raftery \(2007\)](#). This makes proper scoring rules ideal for evaluating forecasts. We will limit our review of scoring methods to rules which are proper.

2.1.2 Storage

For a collaborative forecast project where many researchers are involved and many predictions are collected, computer storage may need to be addressed. As an example of required computer storage, the repository for the COVID-19 Forecast Hub predictions contained 85 million predictions as of February 5, 2022 which required more than 11.6 gigabytes of storage! (cite zoltar??) When determining the goals of a forecast project, there should be consideration of the storage required for different forecast representation types.

2.1.3 Ensemble Modeling

A multi-model ensemble is a statistical model made by combining information from two or more individual statistical models. Private and public decisions are regularly made after combining information from multiple sources. For any given situation, information from one source may provide insight on a subject which other sources fail to capture. Likewise one statistical model may provide insight that another model does not, so that when they are combined in an ensemble the ensemble model outperforms the individual component models.

As probabilistic forecasting becomes more commonplace, so too does ensemble modeling. Multi-model ensembles have been used extensively in weather and climate modeling ([Baran and Lerch, 2018](#), for example), and they have been used increasingly in modeling infectious disease outbreaks see ([Yamana et al., 2016](#), for example). Multi-model ensembles allow for an incorporation of multiple signals - often from differing data sources- and sometimes individual model biases are canceled out or reduced by biases in others ([Reich et al., 2019b](#), see references therein). In several disease outbreak studies, multi-model ensemble forecasts have been shown to outperform individual model forecasts ([Cramer et al., 2021b](#), see references therein).

Construction of an ensemble may be done by combining individual forecast distributions using weighted averages. This has been called stacking [Wolpert \(1992\)](#) or weighted density ensembles [Ray and Reich \(2018\)](#).

2.2 Probabilistic Forecast Representation Types

2.2.1 Parametric distributions

A parametric distribution is a discrete or continuous probability distribution described by a known function $p(x) := p(x|\theta)$ -pmf in the discrete case or pdf in the continuous case. Here θ is a vector of known or unknown parameters contained in a parameter space.

The value of $p(\cdot)$ evaluated at x is defined as $p(x) = P(X = x)$ or the probability that the random variable X takes on x in the space of the random variable. In the continuous case $P(X = x) = 0$ for all x , but the probability that X falls within an interval (a, b) is calculated as

$$P(a < X < b) = \int_a^b p(x)dx \quad (1)$$

Other functions classified by a parametric distribution include the cumulative distribution function (CDF) and the inverse CDF or quantile function. The CDF in the continuous case is defined as

$$F(x) = P(X \leq b) = \int_{-\infty}^x p(t)dt \quad (2)$$

or in the discrete case

$$F(x) = P(X \leq b) = \sum_{k=1}^n p(x_k)dt \quad (3)$$

where x_n is the largest value of X less than or equal to x . The quantile function is defined as

$$Q(p) = \inf\{y \in \mathbb{R} : p \leq F(x)\} \quad (4)$$

returning a quantile value where p is a given probability $0 \leq p \leq 1$.

For a forecast represented as a parametric function with pmf/pdf $p_m(x)$, the accuracy of the forecast may be measured as how likely realized value x^* is to occur. Commonly used proper scoring rules for parametric distributions include the Logarithmic score (LogS), continuous rank probability score (CRPS) [Hersbach \(2000\)](#) [Alves et al. \(2013\)](#) and the interval/Brier score (IS) [Gneiting and Raftery \(2007\)](#) among others. See also [Gneiting and Katzfuss \(2014\)](#) section 3 for more on proper scoring functions. Unless otherwise noted, the following definitions can be found in the review by Krueger.

For a forecast with pdf/pmf $p_m(x)$, the LogS evaluates the probability of the observed value x^* . It is defined as

$$LS(p_m, x^*) = -\log p_m(x^*) \quad (5)$$

The goal is to minimize the LogS, so a forecast $p'_m(x^*)$ is considered superior to $p_m(x^*)$ if $LS(p'_m, x^*) < LS(p_m, x^*)$. The LogS is limited to scoring forecasts with density functions and evaluating those densities only at the point x^* .

The CRPS is a function of a CDF F and so it may be used more extensively than the LogS. For the CDF F_m , the CRPS is defined as

$$CRPS(F_m, x^*) = \int_{-\infty}^{\infty} (F_m(x) - 1_{\{x^* \leq x\}})^2 dy \quad (6)$$

Here too a smaller score indicates a better forecast.

Besides evaluating forecasts, the CRPS may also be used for optimizing weights used to build ensembles under Bayesian model averaging (BMA). Considered the state-of-the-art techniques for combining component distributions into a multi-model ensemble are nonhomogeneous regression (NR), also known as ensemble model output statistics (EMOS), and ensemble BMA, both of which are defined in [Gneiting and Katzfuss \(2014\)](#).

In the context of an ensemble made from component models submitted from various sources, BMA is the better option. In BMA, the final model does not have to be specified beforehand and the resulting forecast will be a discrete mixture distribution of all component forecasts. The general form for an ensemble distribution p^E is

$$p^E(x) = \sum_{m=1}^M w_m p_m(x) \quad (7)$$

where p_m is the m^{th} component forecast distribution and $0 \leq w_m \leq 1$ is a weight assigned to that component where $\sum w_m = 1$. Methods for estimating weights include Maximum Likelihood estimation [Raftery et al. \(2005\)](#), MCMC sampling see [Vrugt et al. \(2008\)](#) and minimizing the ensemble CRPS see [Baran and Lerch \(2018\)](#).

For selecting distribution weights, the CRPS has some nice properties, but can sometimes be difficult to compute. For example, when the forecast is a mixture of a truncated normal distribution (TN) and a truncated lognormal (TL) the CRPS is not available in closed form [Baran and Lerch \(2018\)](#).

Generally computation and evaluation of parametric distributions is not hard. For most commonly used parametric distributions -normal, lognormal, Poisson, gamma, etc.- there is software readily available to compute density, distribution and quantile values. Likewise, requirements for storage are low compared to other representation types that will be discussed. The most common parametric distributions can be fully defined with three or pieces of data including the distribution family and two or three parameters. Table 1 contains enough information to completely define a *Lognormal*(1,0.4) distribution.

family	param1	param2
lnorm	1	0.4

Table 1: Storage example of parametric lognormal distribution with parameters $\mu = 1$ and $\sigma = 0.4$

A completely defined continuous parametric distribution can be evaluated at an infinite number of values. We will call this an infinite resolution.

A drawback of representing a forecast in a parametric distribution is the lack of flexibility in the model selection. Easy computation and evaluation of these models is limited to what is available in software, so certain distributional shapes may be unattainable. The distribution may also assign probability to values outside the range of the forecast. There are remedies for this such as truncation, but these increase the complexity in evaluation and scoring. Requiring a parametric forecast also bars the use of some statistical methods which might be used to create a forecast model including some Bayesian methods where a posterior distribution cannot be computed in closed form.

2.2.2 Sample Distributions

Some forecast model builders may want more flexibility in modeling than a parametric distribution can provide. Methods that require sampling from a posterior distribution or Bootstrap sampling to obtain a sample distribution are examples where parametric distributions may not be appropriate for modeling.

A sample distribution is made of a sample of random variables (X_1, \dots, X_n) where $X_i \sim D$ and D is a distribution. From this, sample statistics such as mean, median, variance and quantiles may be calculated. An empirical cumulative distribution function (ECDF) may also be calculated as

$$F_n(x) = \frac{1}{n} \sum_{i=1}^n \mathbb{I}(X_i \leq x) \quad (8)$$

According to the Strong Law of Large numbers, the sample mean will converge to the expected value of the distribution as $n \rightarrow \infty$ as long as the expected value of that distribution exists. Likewise by the same law the ECDF will converge to the true CDF as $n \rightarrow \infty$. Thus, if a sufficiently large sample is generated from a distribution with an expectation, the sample will closely approximate the true distribution.

For common distribution families it is easy to generate large samples using existing functions in R and other programming platforms. For some distributions for which the mathematical formula is unknown or is not in closed form, more sophisticated methods may be required to generate samples. Bayesian analyses may require a Gibbs or Metropolis-Hastings algorithm, among others, to generate a sample. Such samples are useful in that the true distribution may be closely approximated without knowing the true mathematical form.

Under the sample distribution format, the options that a researcher has for constructing a forecast are more than if they are asked to submit a parametric distribution and the range of characteristics and shapes of distributions is larger. In the last few decades, increased computing power and improvements in MCMC sampling have greatly contributed to growth in the use of predictive distributions [Gneiting and Raftery \(2007\)](#) ([Krueger et al., 2016](#), see examples listed therein).

To properly score a forecast represented by a sample distribution, both the CRPS and LogS may be used. The CRPS has the advantage here of scoring the sample distribution directly since the CDF in [\(6\)](#) may be replaced by the ECDF in [\(8\)](#). To use the LogS to score a forecast, a density function for the sample must be approximated. Common approximations include a kernel density (KD) or Gaussian approximation (GA).

KD is defined by Krueger et. al as

$$\hat{p}_n(x) = \frac{1}{nh_n} \sum_{i=1}^n K\left(\frac{x - X_i}{h_n}\right) \quad (9)$$

where K is a kernel function, and h_n is a suitable bandwidth. GA is defined as

$$\hat{F}_n(x) = \Phi\left(\frac{x - \hat{\mu}_n}{\hat{\sigma}_n}\right) \quad (10)$$

where Φ is the standard normal CDF and $\hat{\mu}_n$ and $\hat{\sigma}_n$ are the empirical mean and standard deviation of the sample (X_i) [Krueger et al. \(2016\)](#).

Krueger et. al compare scoring of MCMC drawn forecast distributions using CRPS and LogS over samples directly using ECDF or over density approximations using KD or GA.

To build an ensemble model from sample distribution forecasts, the BMA construction from (7) may be used only replacing p_m with $\hat{p}_{n_m}^{KD}$ or $\hat{p}_{n_m}^{GA}$ -approximate KD and GA densities. Here as well the optimal weights w_m may be estimated by maximizing the likelihood or minimizing CRPS. Then the ensemble distribution would take the form of a continuous or discrete probability mixture distribution. If the desire is that the ensemble still be a sample distribution, after selecting weights, a sample may be selected by randomly selecting a sample from $(X_n)_m$ with probability w_m .

A potentially large issue with using sample distributions is the amount of storage it would require. For example when making MCMC draws from a posterior distribution, the final sample distribution can have a sample size of thousands or tens of thousands. Maybe not all distributions would require such a large sample size, but sizes of at least dozens or hundreds would be required for each forecast prediction. For any project the size of the CDC flu competition or the COVID-19 Forecast Hub, the storage required would be large and potentially expensive.

2.2.3 Discretized bin distributions

An alternative to parametric distributions and sample distributions which allows for higher flexibility in distribution shape than a parametric distribution and will usually require less storage space than samples is a discretized bin distribution.

A discretized bin probability distribution may be constructed over a set $A = [a, b]$ by partitioning A into a set of K bins $\{B_i\}_{i=1}^K$ where $B_i = [b_{i-1}, b_i)$ and $\cup_{i=1}^K B_i = A$. Based on the problem to be forecasted, researchers will determine the possible range A and select the number of bins and the sizes for each bin. It may be the case that a collaborative project will set the width of all bins to be equal so that $\Delta = b_i - b_{i-1}$ is the same width for all i (McGowan et al., 2019, see for example). To complete the construction, a probability p_i is assigned to each b_i where $\sum_{i=1}^K p_i = 1$. These probabilities are determined by the forecasters. This discrete representation with a given bin and assigned probability is in essence a probability mass function in that the calculation of moments and the the cumulative distribution are done the same as for a discrete parametric distribution. For a random variable X from a binned distribution to calculate the j^{th} moment we may use

$$EX^j = \sum_{i=1}^n \beta_i^j p_i \quad (11)$$

where β_i is a value corresponding to bin B_i (such as $\beta_i = b_i$, $\beta_i = b_{i-1}$ or β_i is the mean or mid value in B_i). The cumulative distribution may be calculated as

$$P(X \leq b) = \sum_{i=1}^n p_i \quad (12)$$

where p_n is the probability for the bin B_n where $b \in B_n$.

If the value to be forecasted takes on discrete values, a common discrete distribution, such as a binary or Poisson distribution, may sometimes be used to assign probabilities to each of the bins. When the values to be forecasted are continuous, a forecaster may need to employ a method of discretization to a forecast distribution. There are a number of possible ways to do this including those outlined by Chakraborty and Subrata Chakraborty (2015).

The CDC has used a discretized bin distribution representation as the format for their annual influenza forecasting competition and other disease outbreak forecast projects. In that context it has become the standard representation [Brooks et al. \(2020\)](#). Much work has been done in evaluating and constructing ensemble models on influenza forecasts represented by discretized bins [McGowan et al. \(2019\)](#); [McAndrew and Reich \(2019\)](#); [Reich et al. \(2019a\)](#).

Because discretized bins can be seen as a pmf, methods for proper scoring already discussed -LogS and CRPS- are useable and BMA is valid method for ensemble construction. Reich et. al used BMA to combine multiple forecasts from the flu competition. They built and compared ensemble models with different weighting schemes including equally weighted components $w_j = 1/J$ and estimating weights according to the model specification. To estimate weights they used the Expectation Maximization (EM) algorithm ([Reich et al., 2019b](#), see supplementary material within for details).

The exact amount of information required for a discretized bin forecast will vary depending on the permitted range of the forecast and the desired resolution. In the CDC flu contest, a forecast might have 131 bins between 0% and 13% -bins having increments of 0.1 or 0.001%- with corresponding probabilities in each. This makes 262 pieces of information per prediction. For any binning scheme of more than two bins, the information requirement for discretized bins will be higher than for parametric distributions. Table 2 illustrates what a discretized Lognormal(1,0.4) distribution truncated over $[0, 8]$ looks like in 41 equally spaced discretized bins. The discretization was done such that

$$p_i = \int_{b_{i-1}}^{b_i} p^{TL}(x)dx \quad (13)$$

where p^{TL} is the pdf of a truncated Lognormal(1,0.4). This is similar to Methodology-IV from Chakraborty [Chakraborty \(2015\)](#); [Kemp \(2004\)](#).

Submitted as a forecast prediction, the distribution in table 2 includes 82 pieces of data. For parts of the CDC influenza competition some forecasts included up to 262. This is far less storage

bin	prob
...	...
[1.4, 1.6)	0.04414
[1.6, 1.8)	0.05896
[1.8, 2.0)	0.07032
[2.0, 2.2)	0.07172
[2.2, 2.4)	0.07955
...	...

Table 2: Storage example of discretized lognormal with $\mu = 1$ and $\sigma = 0.4$

than the possible thousands of draws from a sample distribution but is still much larger than the three or five data pieces required to report a lognormal or truncated lognormal distribution.

Depending on what is known about the problem to be solved, creating the right binning scheme may be a challenge. Because there must be a finite number of bins, forecast distributions often have a finite support. Where the range of possible outcomes to a problem is not well known, the right binning scheme may be hard to produce.

2.2.4 Quantile representation

When deciding how forecasts should be represented in the COVID-19 Forecast Hub, the time pressure of generating forecasts and the consideration of the large range for possible outcomes both contributed to the COVID-19 Forecast Hub decision to forego trying to create the right binning scheme and use sets of prediction intervals to represent uncertainty in predictions [Bracher et al. \(2021\)](#).

The Forecast Hub requires predictions to be submitted as 11 or three nominal intervals -depending on the specific target and unit to forecast- and a median. Using this interval or quantile representation leads to a number of changes in the scoring and ensemble building of forecasts.

A quantile forecast is constructed as follows. For N given quantiles $\alpha_1, \dots, \alpha_N$; q_1, \dots, q_N are the values such that

$$P(Y \leq q_1) = \alpha_1, P(Y \leq q_2) = \alpha_2, \dots, P(Y \leq q_N) = \alpha_N \quad (14)$$

When the quantiles are reported as prediction intervals we have

$$P(Y \leq q_1) = \alpha_1, P(Y \leq q_2) = \alpha_2, \dots, P(Y \leq q_{N-1}) = 1 - \alpha_2, P(Y \leq q_N) = 1 - \alpha_1 \quad (15)$$

What is possible for scoring forecasts or building ensemble models in other representations is often not possible with quantile forecasts. The shape of a distribution is also not known and in fact, nothing is known about the tails or the uncertainty beyond the most extreme reported quantile values. In the Forecast Hub predictions, nothing is reported about the range below the 1st quantile or above the 99th. Yet the quantile representation has its advantages.

Quantile representation allows for forecasters to submit fairly detailed forecasts without restricting the range of possible values. Since quantiles are easily calculated from any regular distribution type -e.g. quantile function for parametric functions or calculating sample quantiles- we consider quantile forecasts to be highly flexible in terms of what methods forecasters may employ in modeling.

To score predictions from a quantile or prediction interval format, neither the LogS nor the CRPS may be used, but another proper scoring rule the IS may be used. For an observed outcome x^* and a prediction interval (l, r) where l and r are the $\alpha/2$ and $(1 - \alpha/2)$ quantiles that bound the central $(1 - \alpha)$ prediction interval the IS is defined as

$$IS_\alpha(l, r; x^*) = (r - l) + \frac{2}{\alpha}(l - x^*)1\{x^* < l\} + \frac{2}{\alpha}(y - r)1\{x^* > r\} \quad (16)$$

This is a sum -weighted by α - of the width of the interval and the distance between x^* and the interval (if x^* is not captured in the interval) [Gneiting and Katzfuss \(2014\)](#). The IS requires only a single central $(1 - \alpha) \times 100$ prediction interval.

When an interval forecast is made up of multiple intervals each with different α levels, the weighted interval score (WIS) may be evaluated. Bracher et. al use the WIS to evaluate COVID-19 interval forecasts [Bracher et al. \(2021\)](#). There are multiple versions of the WIS -some

of which are described in Bracher et. al- but the version used by the COVID-19 Forecast Hub for a forecast of K intervals is defined as follows

$$WIS_{0,K}(F_m, y^*) = \frac{1}{K + 1/2} \times (w_0 \times |y^* - med| + \sum_{k=1}^K \{w_k \times IS_{\alpha_k}(F_m, y)\}) \quad (17)$$

where $w_k = \alpha_k/2$ is the weight on the k^{th} interval. With that choice of weights, it may be shown that the WIS approximates the CRPS (Bracher et al., 2021, see S1 Text therein).

Bogner, Liechti and Zappa compared scoring forecasts of quantiles with the Quantile Score (QS) similar to the interval score and scoring distribution functions fit to those quantiles using the CRPS Bogner et al. (2017). The CRPS corresponds to the integral of the QS over all possible thresholds rather than just specific quantiles, so it more effectively reveals deficiencies in parts of the distribution and especially in the tails past the end points of quantiles used in QS or IS. Thus there may be something lost in terms of scoring when the WIS is used.

Like the CRPS, not only does the WIS provide an easily interpretable proper score for interval forecasts, but it may also be useful when building an ensemble forecast. The ensemble forecast constructed by the COVID-19 Forecast Hub was made as an equally-weighted average of forecasts from the component models. More specifically, each quantile value of the ensemble was the average of values from all models corresponding to the same quantile Ray et al. (2020). For M models each with K quantiles, the k^{th} ensemble quantile q_k^E is calculated as

$$q_k^E = \sum_{m=1}^M \nu_m q_k^m \quad (18)$$

where ν_m is the weight assigned to each forecast and $\sum \nu_m = 1$. In the Forecast Hub model, $\nu_m = \nu = 1/M$. Where the overall mean or a weighted mean may be used, the median may also be used. Brooks et. al compare performance of the COVID-19 ensemble between equally-weighted means, weighted means and median value constructions. Brooks et al. (2020) (is is appropriate to cite this blog??) In their report, they show that weighted means and median constructions tend to outperform an equally-weighted mean construction. To come up with optimal weights, they select values ν_m from (18) which minimize the WIS of the ensemble forecast.

Averaging quantiles in this way is the same as quantile averaging or Vincentization only with an incomplete quantile function. It thus carries with it many of the same characteristics. Quantile averaging or Vincentization for a complete distribution is defined as

$$F_v^{-1}(\alpha) = \sum_{m=1}^M w_m F_m^{-1}(\alpha) \quad (19)$$

where $F_m^{-1}(\alpha) = \inf\{y : F_m(y) \geq \alpha\}$ for $0 < \alpha \leq 1$. Some notable characteristics are that it is more likely for the ensemble distribution to be unimodal than it is under linear averaging of densities like BMA [Busetti \(2017\)](#) Under some circumstances, such as when member distributions are exponential, Weibull or logisitic the aggregated distribution is the same [Ratcliff \(1979\)](#). It produces smoother distributions than BMA according to Schepen and Wang [Schepen and Wang \(2015\)](#). Lichtendahl et. al conclude that quantile averaging produces sharper forecasts and tends to perform better in scoring than probability averaging [Lichtendahl Jr et al. \(2013\)](#).

As in sample distributions and binned distributions, data storage for interval forecasts will depend on the desired clarity of resolution. For the COVID-19 forecasts submitted to the Forecast Hub, 23 quantile values are requested for the quantiles (0.01, 0.025, 0.05, 0.10, \dots , 0.95, 0.975, 0.99). This included a median along with eleven confidence intervals [Bracher et al. \(2021\)](#). This requires forecasters to submit 46 values in each short-term forecast (some of the longer term forecasts only include seven quantiles). In terms of storage, this is an improvement over requirements for the flu contest. Table 3 shows what a submission of 23 quantiles from a *Lognormal*(1,0.4) looks like.

quantile	0.01	0.025	0.05	...	0.95	0.975	0.99
value	1.07137	1.2404	1.40689	...	5.18328	5.82391	6.58783

Table 3: Storage example of quantiles and values of lognormal with $\mu = 1$ and $\sigma = 0.4$

2.3 Discrete mixture of parametric distributions

A discrete mixture of parametric distributions forecast representation may be an attractive alternative to the four types already discussed. A discrete mixture would allow for many

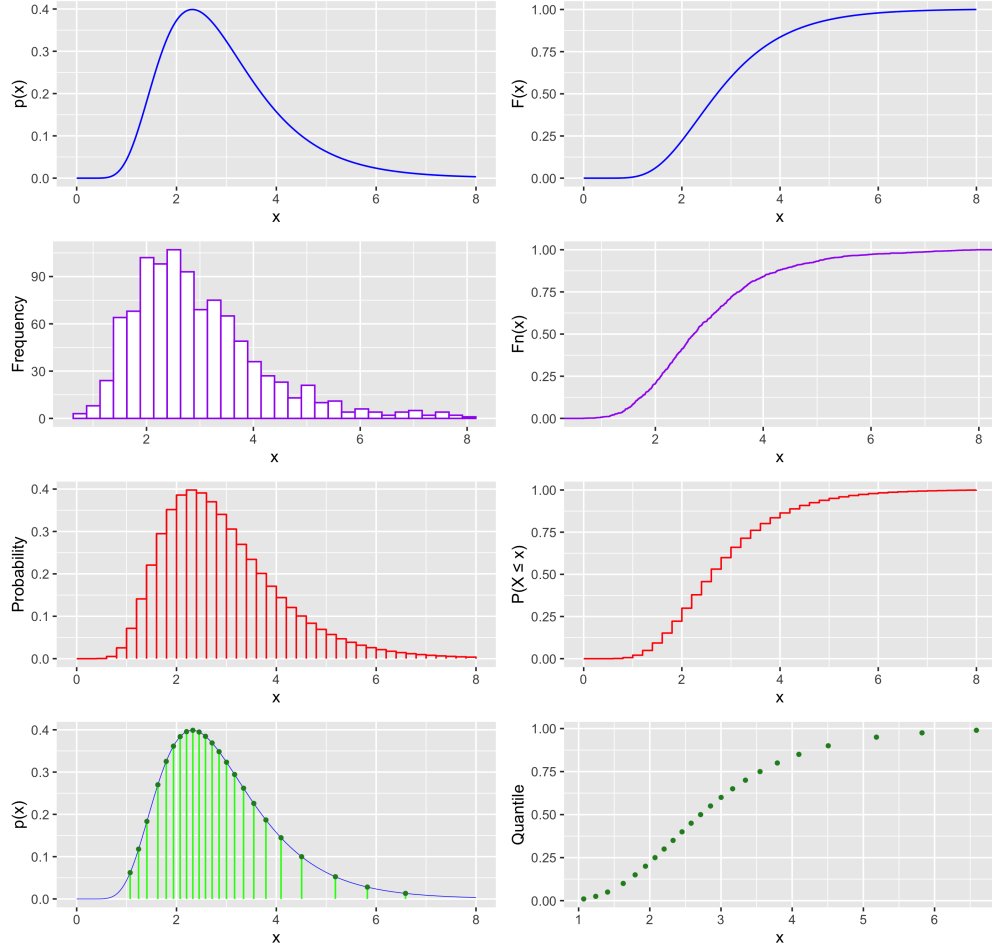


Figure 2: This figure compares the densities and CDFs of forecast representation types discussed in the left and right columns respectively. Each is generated from a $\text{lognorm}(1,0.4)$ distribution. Blue shows the density and CDF functions. Purple shows a histogram and ECDF of 1,000 samples. Red shows bin probabilities and the CDF function. Green shows quantiles/intervals with corresponding values.

distribution shapes, a high resolution, storage comparable to that of discretized bins and quantile representations and the use of the more common methods for scoring and ensemble building.

A discrete mixture distribution may be constructed in the same way as the ensemble described in section 2.2.1 in equation (7) where for T distributions with pdfs $p_t(x)$ and $z_t > 0$ and $\sum_{t=1}^T z_t = 1$

$$p^M = \sum_{t=1}^T z_t p_t(x) \quad (20)$$

Like the standard parametric model, a discrete mixture may be evaluated with existing software (like `distr` R package [Camphausen et al. \(2007\)](#)). And scoring may be done by using the LogS, CRPS and IS. The resolution of a discrete mixture is infinite and the support need not be limited by submission requirements.

A mixture distribution may be much more flexible than a single component parametric distribution in terms of distributional shape. According to McLachlan and Peel, a finite mixture of normal densities with common variance can be used to approximate arbitrarily well any continuous distribution [Peel and MacLachlan \(2000\)](#) (see also [Nguyen and McLachlan \(2019\)](#)). Thus, for an unconventional probability distribution -such as MCMC posterior samples- it may be reasonable to approximate the distribution by constructing a mixture of normals.

An ensemble model may be built by using (7) only replacing p_m with p_m^M in 20. Solving for weights may also be done similarly, though with the added complexity of component models being mixture distributions the computation is likely to be more expensive. An example where this may be especially true is when minimizing of CRPS, but the exact mixture doesn't result in CRPS in closed form [Baran and Lerch \(2018\)](#). In large projects like the COVID-19 Forecast Hub, if an equal weight is not assigned to each component, it may be determined that models not reaching a certain standard of predictive performance are assigned an ensemble weight of 0. This would simplify an ensemble model to include only the best performing forecasts.

Depending on the number of components a forecaster includes in a mixture forecast, the amount of storage per forecast might be as little as for a parametric forecast and as much as a

Representation	Scoring	Flexibility/Information	Ensemble	Storage Requirement
Parametric	LogS, CRPS, IS	Limited to common distribution families. Infinite resolution	BMA	Low 3-6 values per prediction
Samples	CRPS and IS. LogS after smoothing	Any shape. Resolution may be very high but depends on sample size	BMA after smoothing. Resampling otherwise	Hundreds or thousands of values per prediction
Discretized bins	LogS, CRPS, IS	Any shape allowed, but limited by binning scheme. Range also may be limited	BMA	Depends on binning scheme but dozens to hundreds of values
Quantiles	IS, WIS	Shape unknown but with enough quantiles there is still a decent amount of information. No tail information	Quantile averaging	Depends on requested quantiles. Perhaps dozens of values
Discrete Mixture	LogS, CRPS, IS. May be more limited by computation	With enough component distributions may approximate any distribution shape. Infinite resolution	BMA	3 values to dozens per forecast depending on number of components

Table 4: This table compares scoring, information, ensemble building and storage requirements for the different forecast representations discussed.

forecast hub will allow. In section 3 we show how a mixture distribution may be constructed, scored and used to construct an ensemble using software available in R.

Table 4 shows how a discrete mixture distribution forecast compares with the other formats discussed in terms of methods for scoring, information and resolution provided, methods for ensemble building and computer storage requirement.

3 Mixture distributions in a collaborative forecast project

The CDC and the COVID-19 Forecast Hub as well as other collaborative projects have their own established systems for receiving, evaluating and constructing ensemble forecasts. A transition from using binned distributions or quantile forecasts to using discrete mixture distributions would require a few adjustments to those systems. In this section we outline how some of these adjustments may be implemented as well as present some tools which may be used to evaluate forecasts and construct ensemble forecasts from multiple models.

3.1 Submission format

For a collaborative forecast project to run smoothly, model submissions from all modelers should all follow the same format. For both the CDC flu forecasting and the COVID-19 Forecast Hub, teams are provided a .csv spreadsheet with columns similar to those in tables 2 and 3. Additionally columns for Location, Target, Type and Unit are included which serve as indicators of the specific prediction. The values for all variables except Value come from a specified list from the forecast center. Value is the probability assigned to a bin when the Type is Bin.

Location	Target	Type	Unit	Bin	Value
US National	Season Onset	Point	week	NA	.
US National	Season Onset	Bin	week	0.0	.
US National	Season Onset	Bin	week	0.1	.
...

Table 5: Example of a submission file for an discretized forecast like those in the CDC influenza forecast competition. To match a quantile submission format as used by the COVID-19 Forecast Hub, possible values for the variable Type are Point and Quantile, the Bin column is replaced with a Quantile column and Value is the value of the specified quantile in the forecast.

For the quantile forecasts of the Forecast Hub, the main changes to table ?? are that the possible values for Type are Point and Quantile, and rather than using a Bin variable Quantile is used with values specified by the Hub. Then Value rather than a probability is the forecasted value for the specified quantile. Per prediction there may be a couple dozen rows, up to 24 in the Forecast Hub, or over 100 rows as in some flu forecasts.

In the discrete mixture format, an adjustment may be made to the format of table 6 where each row represents a component in a mixture distribution. The variables Bin/Quantile and Value are removed and replaced with Family, Param1, Param2 and Weight where Family is the distribution family, Params are the parameters for the component distribution and Weight is the weight w_i for that component.

Location	Target	Type	Unit	Family	Param1	Param2	Weight
US National	Season Onset	Dist	week	norm	μ_n	σ_n	.
US National	Season Onset	Dist	week	lnorm	μ_l	σ_l	.
...

Table 6: Example of a submission file for a disease outbreak forecast using mixture distribution representation

A forecast center may want to limit the number of components allowed per prediction. For reference, a mixture distribution prediction following the format in table 6 with 18 components would require $18 \times 8 = 144$ cells submitted. This is the same number of cells submitted for a Forecast Hub forecast with 23 quantiles and a point prediction according to the format in table 6.

Throughout the remainder of this section, explanations of how to work with mixture distributions along with R code which will be given to demonstrate constructing a mixture distribution from a forecast submission, scoring the forecast and building an ensemble forecast.

3.2 Mixture construction and scoring tools

From a submitted forecast, a prediction may be found by Location, Target and Unit (or whatever other indication variables are used by collaborators) and evaluated. We found the `distr` package [Camphausen et al. \(2007\)](#) useful for constructing a mixture distribution given information like that from table 6.

The function `UnivarMixingDistribution` takes as arguments a list of distributions and a vector of weights for each distribution and a object of class `AbscontDistribution` is returned. The object is in essence a mixture distribution of the form of (20). Functions for density, distribution and quantile functions as well as for random sampling may then be used on the

mixture distribution object for evaluation. Using those functions and the created object, LogS and CRPS are easily calculable.

We wrote a function **MakeDist** (see Appendix) which takes on a data frame with variables Family, Param1, Param2, Param3 and Weight and returns a mixture distribution of component distributions specified in the rows.

Beginning with two forecast submissions of the form of table 6, say we select from each a prediction of an event for the same location, target and unit. The following two data frames `preddf1` and `preddf2` are the resulting predictions.

family	param1	param2	param3	weight
Lnorm	2	1	NA	0.3
Norm	2.1	1	NA	0.7

Table 7: Toy prediction 1

family	param1	param2	param3	weight
Norm	1.5	1	NA	0.4
Norm	4	2	NA	0.6

Table 8: Toy prediction 2

The code here shows these two data frames and how the **MakeDist** function is used to create the distributions in R. Once the distributions are created as **AbscontDistribution** classes, then functions for evaluating a pdf and a CDF for each are created.

```
> preddf1
> # family param1 param2 param3 weights
> #1  Lnorm    2.0      1     NA     0.3
> #2   Norm    2.1      1     NA     0.7
>
> preddf2
> # family param1 param2 param3 weights
> #1   Norm    1.5      1     NA     0.4
```

```

> #2    Norm    4.0      2    NA    0.6
>
>
>
> #make mixture distributions from prediction submissions
> mdist1 <- MakeDist(preddf1)
> mdist2 <- MakeDist(preddf2)
> #make pdfs for mixture predictions
> dmdist1 <- function(x) {distr::d(mdistr1)(x)}
> dmdist2 <- function(x) {distr::d(mdistr2)(x)}
> #make cdfs for mixture predictions
> pmdist1 <- function(x) {distr::p(mdistr1)(x)}
> pmdist2 <- function(x) {distr::p(mdistr2)(x)}

```

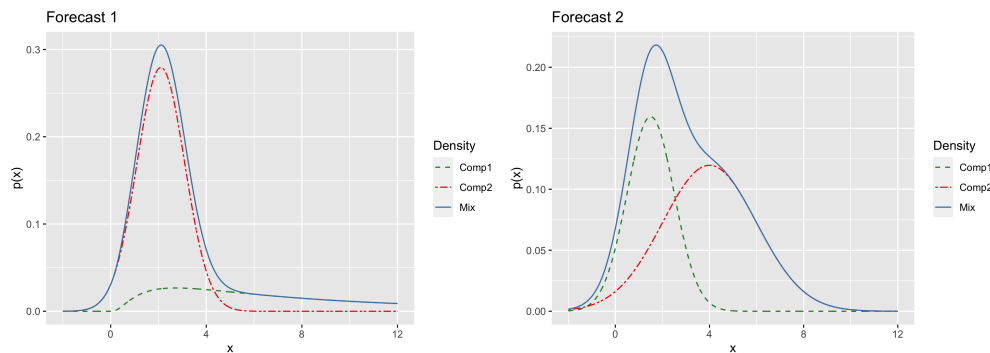


Figure 3: These plots show the density functions of the two toy mixture distribution forecasts along with the component density functions scaled by the corresponding weights

The LogS or the CRPS may then be calculated for each forecast using the pdf and cdf functions respectively. Here we will assume that the realized value which both forecasts attempted to predict was 3. The `CRPS()` function here is included in the appendix of this manuscript. It can be seen here that by the LogS, the forecast from table 7 outperforms that from table 8. Yet under the CRPS the performance is the other way around.

```

> #realized observation
> ob <- 3
> #LogS for predictions at the realized observation
> -log(dmdist1(ob)) # => 1.547238
> -log(dmdist2(ob)) # => 1.848796
> #CRPS for predictions at the realized observation
> CRPS(pmdist1,y=ob) # => 0.6348212
> CRPS(pmdist2,y=ob) # => 0.5306083

```

3.3 Ensemble construction

To construct an ensemble distribution from multiple mixture distributions, the `UnivarMixingDistribution` is again a useful tool. The function takes two or more `AbscontDistribution` or mixture distributions and a vector of weights corresponding to each mixture. A new `AbscontDistribution` object is returned as an ensemble of mixture distributions see (7).

At the onset of a collaborative forecast before the various models can be scored based on true observations, it may make sense to assign equal weight to each component distribution in an ensemble. As a project progresses however, assigning weights based on past performance may be desired. As mentioned in section 2.2.1, weights may be selected by maximizing the likelihood of (7) or by minimizing the CRPS. Another method of selecting weights is to use the posterior model probability.

If we have T models (M_t) the posterior model probability of M_t is defined as

$$p(M_t|x) = \frac{p(x|M_t)p(M_t)}{p(x)} = \frac{p(x|M_t)p(M_t)}{\sum_{k=1}^T p(x|M_k)p(M_k)} \quad (21)$$

where $p(\cdot|M_t) := p_m(\cdot)$ is the density function of the model and $p(M_j)$ is the prior probability assigned to the model. Where the prior probabilities for each model are equal or $p(M_t) = 1/T$ for all t (21) is reduced to

$$p(M_t|x) = \frac{p(x|M_t)}{\sum_{k=1}^T p(x|M_k)} \quad (22)$$

The posterior model probability $p(M_t|x^*)$ for an already observed event x^* are then used as weights for a yet unrealized event being forecasted i.e. $w_t := p(M_t|x^*)$ see (7).

From the same toy problem in the previous section, the following code shows how to use the posterior model probability to select weights and construct an ensemble distribution and score the ensemble forecast. The ensemble distribution with components is show in figure 4.

```
> #posterior model probability for calculating weights
> w1 <- pmdist1(ob)/(pmdist1(ob) + pmdist2(ob))
> w2 <- 1-w1
> #w1 => 0.5286434
> #w2 => 0.4713566
>
> #build ensemble with calculated weights
> ensdist <- distr::UnivarMixingDistribution(mdist1,
+                                           mdist2,
+                                           mixCoeff = c(w1,w2))
> #pdf and cdf for ensemble
> densdist <- function(x) {(distr::d(ensdist)(x))}
> pensdist <- function(x) {(distr::p(ensdist)(x))}
> #LogS for predictions at the realized observation
> -log(densdist(ob)) # => 1.678156
> #CRPS for predictions at the realized observation
> CRPS(pensdist,y=ob) # => 0.5486368
```




Figure 4: Toy ensemble forecast with component mixture distributions shown

4 Retrospective Analysis

For large collaborative forecast projects having already established the representation formats for forecasting, it may be difficult for individual forecast teams to adjust to a discrete mixture distribution format. Exact methods for modeling and forecasting may not cater well to submitting a forecast as a discrete parametric mixture distribution.

In this section we attempt to assess whether or not any forecasts from the CDC influenza forecast competition and the COVID-19 Forecast Hub were generated from common continuous parametric distributions. The forecasts we assess are in the forms of discretized bin predictions or sets of predictive intervals. In these formats, not knowing anything about distributions or samples, formal statistical methods for fitting to parametric distributions and assessing fit do not exist. Thus any results stated here may not be stated in terms of statistical certainty.

4.1 CDC Influenza like Illness

The CDC Retrospective Forecasts project at zotrdata.com (cite??) contains over 850,000 probabilistic influenza-like illness predictions for all combinations of eleven regions in the United States and seven targets from twenty-seven different models. These include predictions during

Influenza season between October 2010 and December 2018. All predictions are in the form of discretized probability bins.

We wanted to determine whether or not some or any of these binned forecasts were generated from known uniform, normal, lognormal or gamma continuous distribution families. The process we followed for making such a claim was to fit a distribution to the probabilities provided in the forecast by minimizing

$$\sum_{i=1}^K (p_i - [F(b_i; \theta) - F(b_{i-1}; \theta)])^2 \quad (23)$$

Here $F(\cdot; \hat{\theta})$ is a CDF, p_i is the reported probability for the i^{th} bin $B_i := [b_{i-1}, b_i)$ and K is the number of bins. The fitted parameter vector $\hat{\theta}$ is the solution to

$$\arg \min_{\theta} \sum_{i=1}^K (p_i - [F(b_i; \theta) - F(b_{i-1}; \theta)])^2 \quad (24)$$

If a well known continuous distribution was fit to the submitted binned prediction and the mean sum of squared differences (MSD) or ?? fell below a specified cutoff value, we considered that the binned predictions were plausibly discretized from the continuous distribution.

To determine what values to use as cutoffs we conducted a study where binned distributions were discretized from known continuous distributions, parameters were fit to the binned distributions and the MSD was calculated.

For each of uniform, truncated normal (TN), truncated lognormal (TL) and truncated gamma (TG) distributions the following was done 1,000 times. A set of 131 bins with interval lengths of 0.1 between 0 and 13.1 was constructed. A value μ was selected from a $Uniform(a, b)$ distribution and σ from a $Uniform(.05, 1.6)$ distribution. μ and σ were taken as respectively the center and scale parameters of a TN distribution. For uniform, TL and TG family distributions, model parameters were solved for so that μ and σ were roughly the mean and standard deviation. The minimization was done using the `optim` function in the **R stats** package. Had we accounted for the changes in moment values from the truncation, the means and standard deviations could

have been exactly μ and σ . But we considered the differences small enough that it didn't matter for our purpose.

With a known distribution, probabilities $p_i = F(b_i; \theta) - F(b_{i-1})$ were calculated for each bin. A uniform or truncated distribution from the same family was fit to the bins by minimizing ?? with the resulting distribution function $F(\cdot; \hat{\theta})$. From the fitted distribution $\hat{p}_i = F(b_i; \hat{\theta}) - F(b_{i-1}; \hat{\theta})$ was computed for each of the 131 bins. Finally the mean sum of squared probability differences $\frac{1}{K} \sum_{i=1}^K (\hat{p}_i - p_i)^2$ was calculated. Table 9 below shows MSD value for which 95% of all 1,000 MSDs fell below. Those values were selected as the cutoff values for declaring whether or not a binned distribution was discretized from a common continuous distribution.

Distribution	MSD 95%
Uniform	8.809460e-05
Truncated Normal	1.126683e-07
Truncated Lognormal	3.877829e-06
Truncated Gamma	1.500152e-06

Table 9: If a binned distribution is fit to a parametric distribution and the MSD is smaller than the corresponding cutoff value listed here, we consider that binned distribution to have been discretized from the fit parametric distribution.

Of the 869,638 predictions from the CDC Retrospective Forecast project, we fit each of Uniform, TN, TL and TG distributions to 11,715 of the individual predictions. The MSD between the prediction probabilities and the fit discretized probabilities was calculated. For each prediction, the fitted distribution with the lowest MSD was considered the best fit and if the MSD fell below the corresponding cutoff value listed in 9 we considered that the prediction was discretized from the fit distribution.

Of the 11,715 binned distributions fit to continuous parametric distributions, 2,502 of those fits produced MSD values below the values listed in table 9. Thus we conclude that a proportion of 0.214 of those forecasts were generated from a common parametric distribution. The results for the 11,715 fit distributions, which families the binned distributions were best fit to and whether or not the fits produce an MSD below the cutoff value, are seen in table 10. Figure 5 is an example of the best fit distributions to the same binned probability distribution.

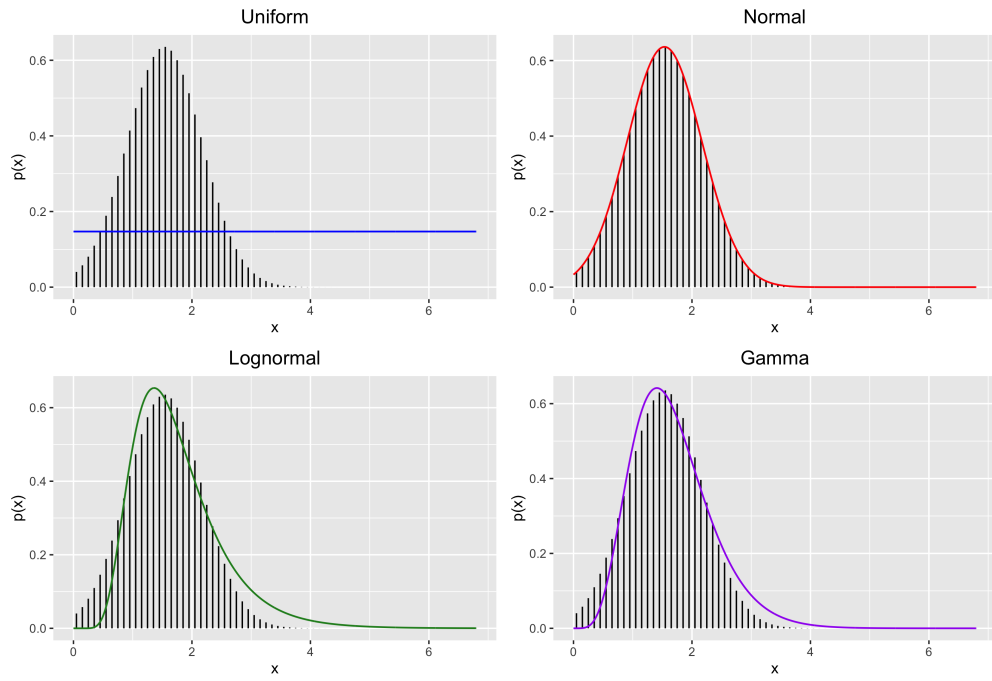


Figure 5: This figure shows the plots of fits from each a uniform, truncated normal, truncated lognormal and truncated gamma distributions to the same binned probability distribution. In this case, the fit to the normal distribution produced an MSD below the normal cutoff value, so we conclude that the binned distribution may have been discretized from truncated normal.

Distribution Family	Total	Total below MSD cutoff value	Proportion from parametric distribution
Uniform	896	669	0.747
Truncated Normal	3,804	198	0.052
Truncated Lognormal	4,501	1,340	0.289
Truncated Gamma	2,514	295	0.117

Table 10: Results for CDC influenza forecast retro analysis

4.2 COVID-19 Forecast Hub

As of January 24, 2022 there were nearly 85,000,000 prediction distributions from 114 different models submitted to the COVID-19 Forecast Hub (cite??). These distributions covered all combinations of 3,202 municipalities (mostly counties) in the United States with 441 targets. The first of these forecasts was submitted in March of 2020 shortly after the initial outbreak of the virus in the US and forecasts had been received weekly since. The forecast representation for these predictions is the quantile representation where a point estimated is submitted along with a number of confidence intervals -three or 11 for most targets. So the predictions typically include seven or 23 quantiles.

To assess whether or not a quantile based forecast was calculated from a well known continuous distribution, we minimized the mean square error (MSE) between the CDF and the given quantiles. The estimated parameters are the solution to $\operatorname{argmin}_{\theta} \sum_{i=1}^m (q_i - F(y; \theta))$ where q_i is the i^{th} given quantile value and $F(y; \theta)$ is the CDF with parameter θ . This least squares estimating is a more common way to fit a model but other methods have been presented including Bayesian Quantile Matching [Nirwan and Bertschinger \(2020\)](#), step interpolation with exponential tails [Quinonero-Candela et al. \(2005\)](#) and the Method of Simulated Quantiles [Dominicy and Veredas \(2013\)](#).

For each prediction used, we fit the quantiles to a uniform, normal, lognormal, gamma and location-scale t CDF. We included the t distribution in case of any symmetric forecasts with heavier tails than in a normal distribution.

If the MSE between a given forecast and a fit CDF fell below a certain cutoff, we considered the quantiles as approximately coming from the fit CDF. To determine the cutoff values, for each of the five distribution families the same procedure was followed 1,000 times.

A decision was randomly made on creating a quantile distribution with seven quantiles or 23 quantiles with probabilities 1/3 and 2/3 respectively. This was done because in the COVID-19 Forecast Hub certain targets require seven quantiles and others require 23. After that decision was made, a random value μ was drawn from a *Uniform*(2,000, 25,000) distribution and another value σ was drawn from a *Uniform*(3,200) distribution. These values were taken as the mean and standard deviation and for each of the distribution families considered, the proper transformations computed to find model parameters corresponding to the distribution. Quantile values were calculated for each quantile. When using a t distribution, a value for degrees of freedom was drawn from a *Uniform*(2,35) distribution. A CDF was fit by minimizing $\sum_{i=1}^m (q_i - F(y; \theta))$ over the parameter vector θ and the MSE value $\sum_{i=1}^m (q_i - F(y; \hat{\theta}))$ was calculated. The MSE value for which 95% of the 1,000 simulated distributions fell below was considered the cutoff and is seen in table 11.

Distribution	Mean Square Error 95%
Uniform	7.220667e-07
Normal	3.097645e-05
Lognormal	1.166775e-07
Gamma	2.953370e-05
Location-scale T	0.4894471

Table 11: If a set of quantiles is fit to a parametric distribution CDF and the MSE is smaller than the corresponding cutoff value listed here, we consider the quantiles to have been calculated parametric distribution.

With these cutoff values selected, models were then fit to quantile forecasts from the COVID-19 Forecasts on zoltardata.com. This was a computationally more difficult problem than fitting distributions to the binned forecasts of the influenza project, so the number of predictions fit is much smaller. The predictions fit were selected from each of the 115 models with model weeks, targets and units selected randomly. In total 2,504 predictions were fit to each of the five

distributions, the MSE was calculated for each and the fit with the lowest MSE was selected as the closest fit. If the MSE fell below the cutoff specified above, we considered that the quantile prediction was approximately from the same distribution as the fit distribution. Of the 2,504 fits, 99 of them produced an MSE value below the corresponding cutoff. Results for fits by distribution are seen in table 12. Figures 6 and 7 show the QQ plots and CDF plots for different distribution fits to the same set of quantiles.

Distribution Family	Total	Total below MSE cutoff	Proportion from parametric distribution
Uniform	239	0	0
Normal	609	74	0.122
Lognormal	694	0	0
Gamma	597	0	0
Location-scale T	365	25	0.068

Table 12: Results for COVID-19 Forecast Hub retro analysis

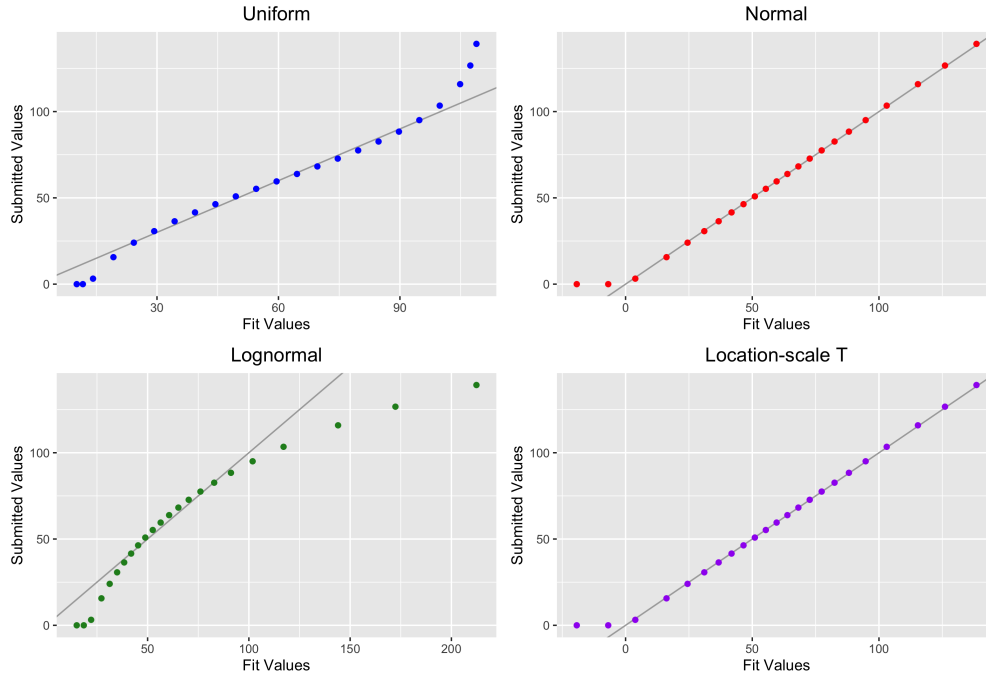


Figure 6: This figure shows qq plots for a set of submitted quantiles against the quantiles the continuous distribution to which it was fit. In no case here was the MSE below the cutoff value of the corresponding distribution.

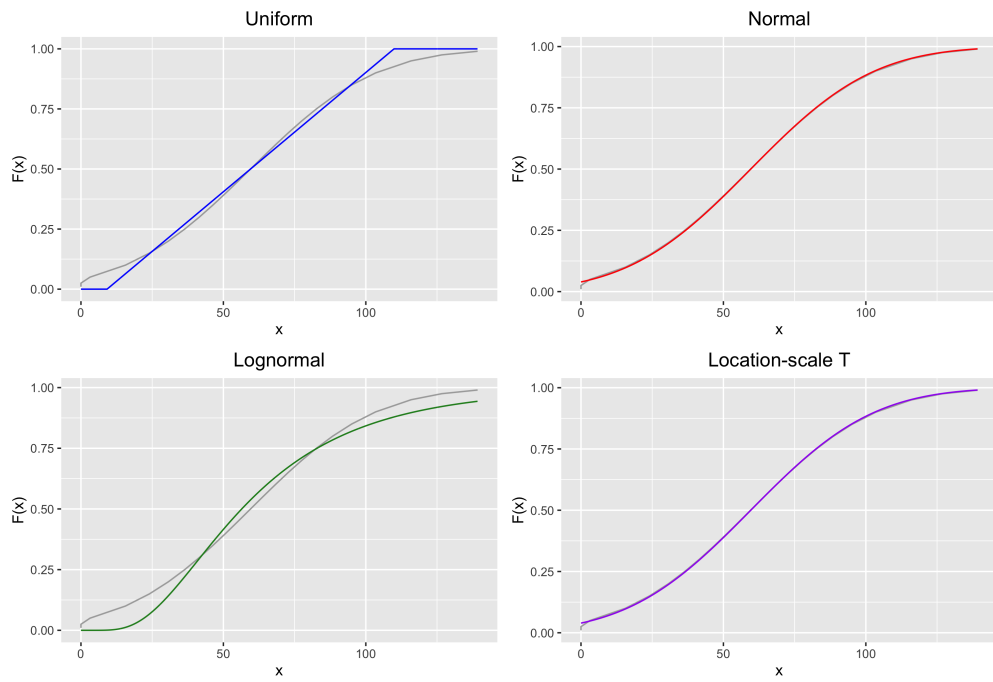


Figure 7: This figure shows the CDF plots for distributions fit to a set of quantiles. The grey line is a line of connected dots from the submitted quantiles and values. In no case here was the MSE below the cutoff value of the corresponding distribution.

The results in the previous two studies suggest that in some cases, forecasters are using methods which produce simple parametric distribution forecasts. Many likely are not. To improve upon these studies, one could fit the binned influenza forecasts or the COVID-19 quantile forecasts to discrete mixture distributions with multiple components. Forecasts that are very close to a mixture distribution may suggest that the forecasters already have models that make transition to a discrete mixture distribution forecast format easy or straightforward.

5 Discussion

In this paper we have reviewed four representation types commonly used in probabilistic forecasting including proper scoring, data storage and ensemble model construction for each type. We presented the discrete mixture distribution representation and we argue that its use in collaborative probabilistic forecasting is preferable to the other representations. In terms of model flexibility, storage and ensemble building it is comparable to discretized bins and interval forecasts but also provides a forecast with an infinite nominal resolution. Thus we advocate its use in future forecasting projects like those done by the CDC and COVID-19 Forecast Hub.

For a number of reasons, the adoption of finite mixture distributions as the submission format may face hurdles. A collaborative forecast center, along with forecasters, using a different representation may simply not want to break from tradition. The implementation of new scoring and ensemble construction methods may be a barrier. One aspect of ensemble construction which received little attention in this paper is the selection of weights for components of an ensemble where each of the components is mixture distribution. Computing requirements could become a concern in such a problem and further research on this may provide ideas of best methods for weight selection.

Another area of recommended research is the use of joint distributions for forecasting. We have only considered here probabilistic forecasting of one event at a time. Number of new infections in one week at one specific location for example. This forecast is presented as a marginal distribution for that specific target, time and location. A joint distribution for forecasting multiple targets, times or locations may sometime be desirable and may require further consideration on how joint mixture distributions could be used as a format in collaboration.

6 References

Alves, J.-H. G., Wittmann, P., Sestak, M., Schauer, J., Stripling, S., Bernier, N. B., McLean, J., Chao, Y., Chawla, A., Tolman, H., et al. (2013). The ncep–fmmoc combined wave ensemble product: Expanding benefits of interagency probabilistic forecasts to the oceanic environment. *Bulletin of the American Meteorological Society*, 94(12):1893–1905.

- Baran, S. and Lerch, S. (2018). Combining predictive distributions for the statistical post-processing of ensemble forecasts. *International Journal of Forecasting*, 34(3):477–496.
- Bogner, K., Liechti, K., and Zappa, M. (2017). Combining quantile forecasts and predictive distributions of streamflows. *Hydrology and Earth System Sciences*, 21(11):5493–5502.
- Bracher, J., Ray, E. L., Gneiting, T., and Reich, N. G. (2021). Evaluating epidemic forecasts in an interval format. *PLoS computational biology*, 17(2):e1008618.
- Brooks, L. C., Ray, E. L., Bien, J., Bracher, J., Rumack, A., Tibshirani, R. J., and Reich, N. G. (2020). Comparing ensemble approaches for short-term probabilistic covid-19 forecasts in the us. *International Institute of Forecasters*.
- Busetti, F. (2017). Quantile aggregation of density forecasts. *Oxford Bulletin of Economics and Statistics*, 79(4):495–512.
- Camphausen, F., Kohl, M., Ruckdeschel, P., Stabla, T., and Ruckdeschel, M. P. (2007). The distr package.
- CDC (2021). Flusight: Flu forecasting.
- Chakraborty, S. (2015). Generating discrete analogues of continuous probability distributions-a survey of methods and constructions. *Journal of Statistical Distributions and Applications*, 2(1):1–30.
- Cramer, E. Y., Huang, Y., Wang, Y., Ray, E. L., Cornell, M., Bracher, J., Brennen, A., Castro Rivadeneira, A. J., Gerding, A., House, K., Jayawardena, D., Kanji, A. H., Khandelwal, A., Le, K., Niemi, J., Stark, A., Shah, A., Wattanachit, N., Zorn, M. W., Reich, N. G., and Consortium, U. C.-. F. H. (2021a). The united states covid-19 forecast hub dataset. *medRxiv*.
- Cramer, E. Y., Lopez, V. K., Niemi, J., George, G. E., Cegan, J. C., Dettwiller, I. D., England, W. P., Farthing, M. W., Hunter, R. H., Lafferty, B., et al. (2021b). Evaluation of individual and ensemble probabilistic forecasts of covid-19 mortality in the us. *medRxiv*.
- Dominicy, Y. and Veredas, D. (2013). The method of simulated quantiles. *Journal of Econometrics*, 172(2):235–247.
- Gneiting, T., Balabdaoui, F., and Raftery, A. E. (2007). Probabilistic forecasts, calibration and sharpness. *Journal of the Royal Statistical Society: Series B (Statistical Methodology)*, 69(2):243–268.
- Gneiting, T. and Katzfuss, M. (2014). Probabilistic forecasting. *Annual Review of Statistics and Its Application*, 1:125–151.

- Gneiting, T. and Raftery, A. E. (2007). Strictly proper scoring rules, prediction, and estimation. *Journal of the American statistical Association*, 102(477):359–378.
- Gneiting, T. and Ranjan, R. (2013). Combining predictive distributions. *Electronic Journal of Statistics*, 7:1747–1782.
- Groen, J. J., Paap, R., and Ravazzolo, F. (2013). Real-time inflation forecasting in a changing world. *Journal of Business & Economic Statistics*, 31(1):29–44.
- Hersbach, H. (2000). Decomposition of the continuous ranked probability score for ensemble prediction systems. *Weather and Forecasting*, 15(5):559–570.
- Kemp, A. W. (2004). Classes of discrete lifetime distributions.
- Krueger, F., Lerch, S., Thorarinsdottir, T. L., and Gneiting, T. (2016). Probabilistic forecasting and comparative model assessment based on markov chain monte carlo output. *arXiv preprint arXiv:1608.06802*, 12.
- Lichtendahl Jr, K. C., Grushka-Cockayne, Y., and Winkler, R. L. (2013). Is it better to average probabilities or quantiles? *Management Science*, 59(7):1594–1611.
- McAndrew, T. and Reich, N. G. (2019). Adaptively stacking ensembles for influenza forecasting with incomplete data. *arXiv preprint arXiv:1908.01675*.
- McGowan, C. J., Biggerstaff, M., Johansson, M., Apfeldorf, K. M., Ben-Nun, M., Brooks, L., Convertino, M., Erraguntla, M., Farrow, D. C., Freeze, J., et al. (2019). Collaborative efforts to forecast seasonal influenza in the united states, 2015–2016. *Scientific reports*, 9(1):1–13.
- Nguyen, H. D. and McLachlan, G. (2019). On approximations via convolution-defined mixture models. *Communications in Statistics-Theory and Methods*, 48(16):3945–3955.
- Nirwan, R.-S. and Bertschinger, N. (2020). Bayesian quantile matching estimation. *arXiv preprint arXiv:2008.06423*.
- Peel, D. and MacLachlan, G. (2000). Finite mixture models. *John & Sons*.
- Quinonero-Candela, J., Rasmussen, C. E., Sinz, F., Bousquet, O., and Schölkopf, B. (2005). Evaluating predictive uncertainty challenge. In *Machine Learning Challenges Workshop*, pages 1–27. Springer.
- Raftery, A. E., Gneiting, T., Balabdaoui, F., and Polakowski, M. (2005). Using bayesian model averaging to calibrate forecast ensembles. *Monthly weather review*, 133(5):1155–1174.
- Ratcliff, R. (1979). Group reaction time distributions and an analysis of distribution statistics. *Psychological bulletin*, 86(3):446.

- Ray, E. L. and Reich, N. G. (2018). Prediction of infectious disease epidemics via weighted density ensembles. *PLoS computational biology*, 14(2):e1005910.
- Ray, E. L., Wattanachit, N., Niemi, J., Kanji, A. H., House, K., Cramer, E. Y., Bracher, J., Zheng, A., Yamana, T. K., Xiong, X., et al. (2020). Ensemble forecasts of coronavirus disease 2019 (covid-19) in the us. *MedRxiv*.
- Reich, N. G., Brooks, L. C., Fox, S. J., Kandula, S., McGowan, C. J., Moore, E., Osthus, D., Ray, E. L., Tushar, A., Yamana, T. K., et al. (2019a). A collaborative multiyear, multimodel assessment of seasonal influenza forecasting in the united states. *Proceedings of the National Academy of Sciences*, 116(8):3146–3154.
- Reich, N. G., McGowan, C. J., Yamana, T. K., Tushar, A., Ray, E. L., Osthus, D., Kandula, S., Brooks, L. C., Crawford-Crudell, W., Gibson, G. C., et al. (2019b). Accuracy of real-time multi-model ensemble forecasts for seasonal influenza in the us. *PLoS computational biology*, 15(11):e1007486.
- Schepen, A. and Wang, Q. (2015). Model averaging methods to merge operational statistical and dynamic seasonal streamflow forecasts in australia. *Water Resources Research*, 51(3):1797–1812.
- Taylor, J. W. (2021). Evaluating quantile-bounded and expectile-bounded interval forecasts. *International Journal of Forecasting*, 37(2):800–811.
- Vrugt, J. A., Diks, C. G., and Clark, M. P. (2008). Ensemble bayesian model averaging using markov chain monte carlo sampling. *Environmental fluid mechanics*, 8(5):579–595.
- Wolpert, D. H. (1992). Stacked generalization. *Neural networks*, 5(2):241–259.
- Yamana, T. K., Kandula, S., and Shaman, J. (2016). Superensemble forecasts of dengue outbreaks. *Journal of The Royal Society Interface*, 13(123):20160410.

7 Appendix

The following is the code to make the `MakeDist()` function introduced in [3.2](#).

```
> MakeDist <- function(distsdf){
+
+   distdf <- distsdf[distsdf[,1] != 'Lst',]
+   tdist <- distsdf[distsdf[,1] == 'Lst',]
+ }
```

```

+ fun_dist <-
+   apply(distdf, FUN=function(x) {
+     paste('distr::',x[1], '(',
+       ifelse(!is.na(x[2]) & (!is.na(x[3]) | !is.na(x[4])),
+         paste(x[2],',',sep=''),
+         ifelse(!is.na(x[2]) & is.na(x[3]) & is.na(x[4]),
+           x[2], '')),
+       ifelse(!is.na(x[3]) & !is.na(x[4]),
+         paste(x[3],',',sep=''),
+         ifelse(!is.na(x[3]) & is.na(x[4]), x[3], '')),
+       ifelse(!is.na(x[4]),x[4], ''), ' '),sep='')
+   }, MARGIN = 1
+ )
+
+ fun_tdist <- apply(tdist, FUN=function(x) {
+   paste0('distr::Td(',x[4],')*',x[3], '+', x[2])
+ }, MARGIN = 1
+ )
+
+ dist_args <- paste(fun_dist, collapse=',',sep='')
+ tdist_args <- paste0(fun_tdist,collapse=',')
+ args <- ifelse(tdist_args!='',paste(dist_args,tdist_args,sep=','),dist_args)
+
+ weights <- c(distdf[,5],tdist[,5])
+ mixString <- paste('distr::UnivarMixingDistribution(',
+   args,',mixCoeff=weights)',sep='')
+ mixDist <- eval(parse(text=mixString))

```

```

+
+   return(mixDist)
+ }

```

The following is the code used to make the function `CRPS()` used in section [3.2](#).

```

> crps_integrand <- function(x,dist,y) {(dist(x) - as.numeric(y <= x))^2}
> CRPS <- function(y,dist) {
+   int <- integrate(crps_integrand,-Inf,Inf,y,dist=dist)
+   return(int$value)
+ }

```